Machine Learning

Lecture 1: Overview

Feng Li

fli@sdu.edu.cn
https://funglee.github.io

School of Computer Science and Technology
Shandong University

Fall 2018

- Instructor: Prof Feng Li
- Office: N3-312-1
- Affiliation: Institute of Intelligent Computing (https://sduiic.github.io)
- Education:
  - 2010-2015, PhD, Nanyang Technological University, Singapore.
  - 2007-2010, MS, Shandong University, China.
  - 2003-2007, BS, Shandong Normal University, China.
- Employment:
  - 2015-Present, Assistant Professor, Shandong University, China.
  - 2014-2015, Research Fellow, National University of Singapore, Singapore.
- Research Interests:
  - Applied Optimization
  - Distributed Algorithms and Systems
  - Wireless Networking
  - Mobile Sensing and Computing

Suggested Readings

- Zhihua Zhou, Machine Learning, Tsinghua Press, 2016
- Tom M. Mitchell, Machine Learning (1st Ed.), China Machine Press, 2008
- Ian Goodfellow, Yoshua Bengio, Deep Learning, People's Posts and Telecommunications Press, 2016 (Online: `http://www.deeplearningbook.org/`)
- Trevor Hastie, The Elements of Statistical Learning: Data Mining, Inference, and Prediction (2nd Ed.), World Publishing Corporation, 2015
- Simon Haykin, McMaster, Neural Networks and Learning Machines (3rd Ed.), China Machine Press, 2009
- Christopher M. Bishop, Pattern Recognition and Machine Learning (1st Ed.), Springer, 2006

Course Information

- Credits
  - 64 hours (4 hours/week $\times$ 16 weeks)
  - Final Exam: Exam week
- Grades
  - Labs/Projects: 25%
  - Homework : 25%
  - Final exam : 50%
- Website: `https://funglee.github.io/ml/ml.html`
- Teaching Assistants (TAs):
  - Ms Yuan Yuan (yuan930126 AT 163 DOT com)
  - Mr Qi Luo (luoqi4110217 AT hotmail AT com)
  - Mr Xunjian Li (xunjianli AT mail DOT sdu DOT edu DOT cn)
  - Ms Ye Feng (vaneness1998 AT gmail DOT com).

Course Requirements

- Linear algebra
- Calculus
- Probability theory
- Statistics
- Information theory
- Convex Optimization

Remarks

- Lectures are important, but not enough
- You should review what have been taught with more hours than the class hours/week
- You should be familiar with all terminologies related with this course
- You should understand the theories behind machine learning techniques
- Practice what you have learned

History of AI

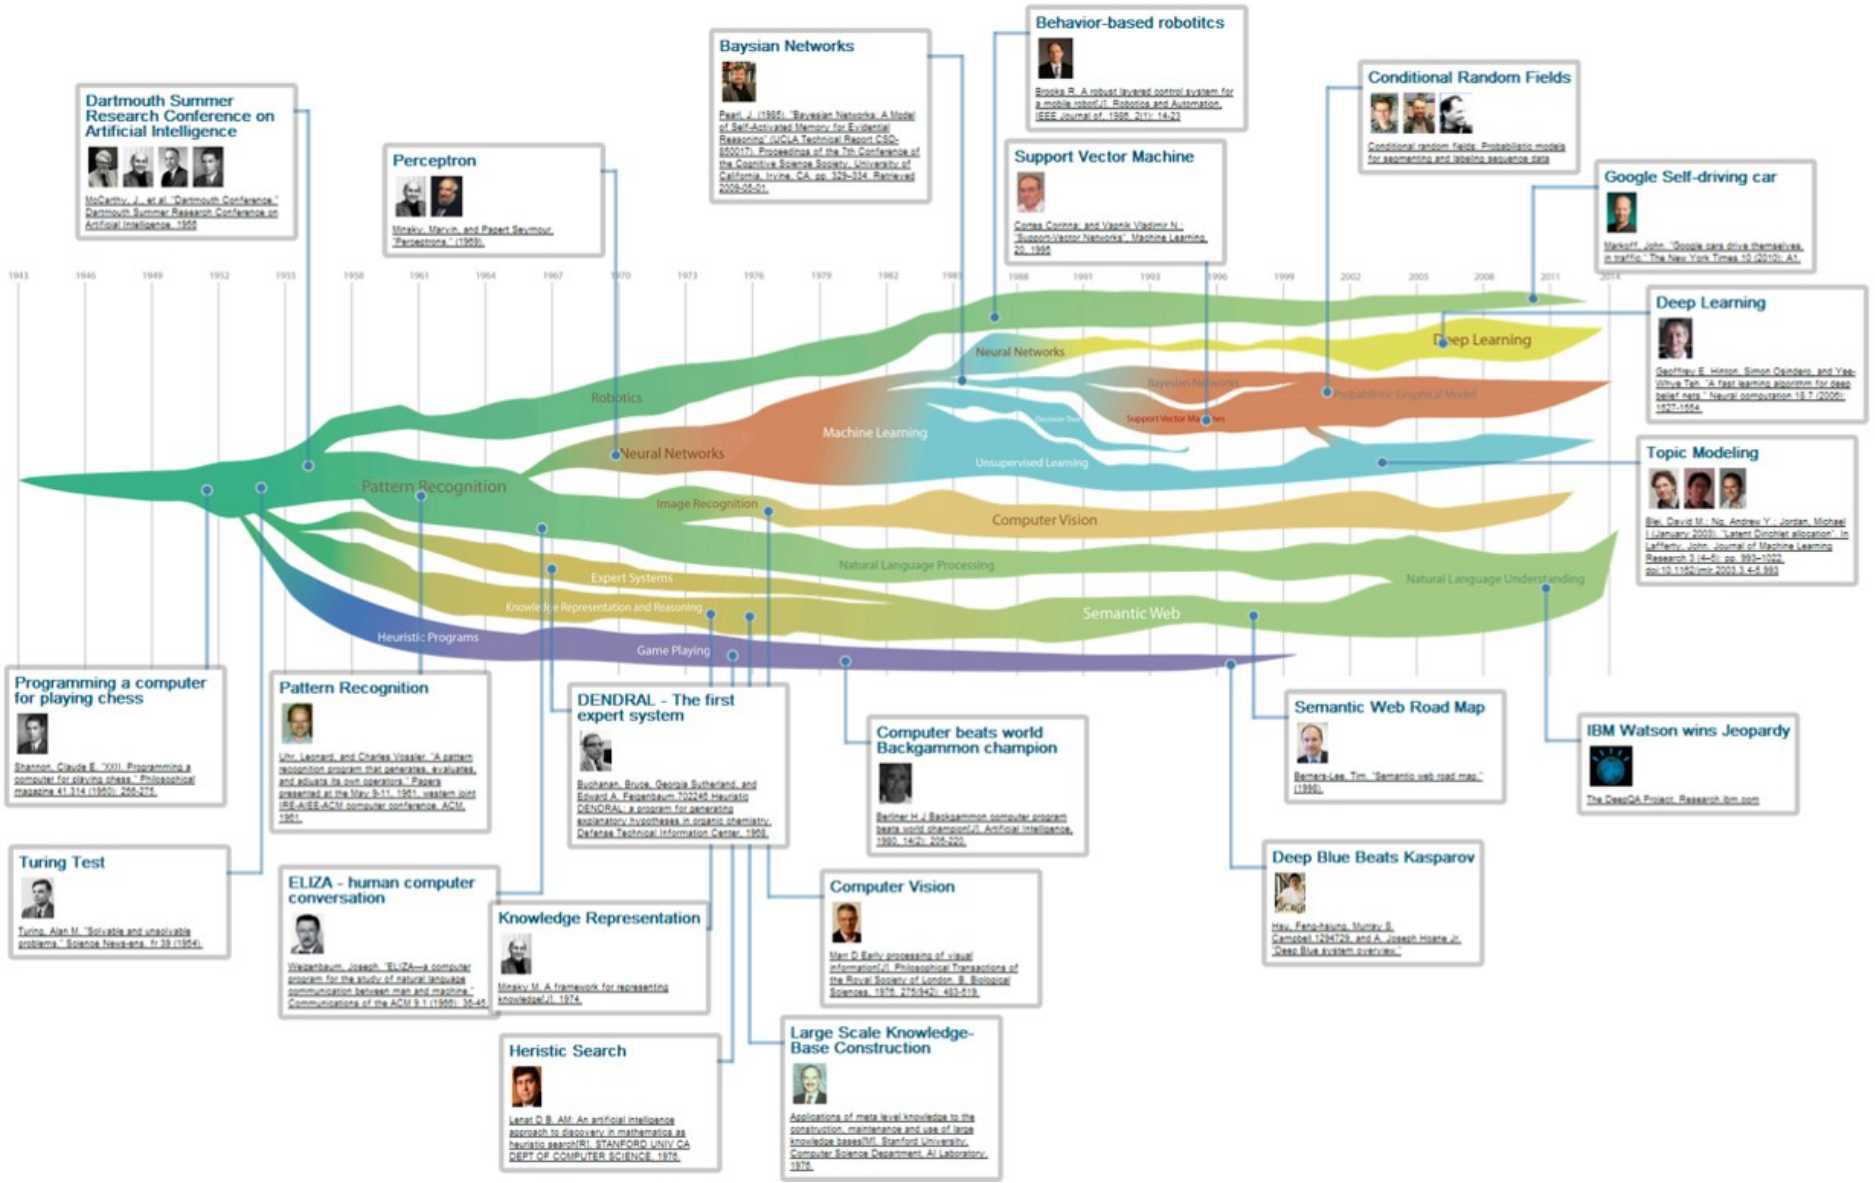

History

- 1950s
  - Turing test
  - Samuel's checker player
- 1960s:
  - Neural networks: Perceptron
  - Pattern recognition
  - Minsky and Papert prove limitations of Perceptron
- 1970s:
  - Symbolic concept induction
  - Winston's arch learner
  - Expert systems and the knowledge acquisition bottleneck
  - Quinlan's ID3
  - Michalski's AQ and soybean diagnosis
  - Scientific discovery with BACON
  - Mathematical discovery with AM

History

- 1980s:
  - Advanced decision tree and rule learning
  - Explanation-based Learning (EBL)
  - Learning and planning and problem solving
  - Utility problem
  - Analogy
  - Cognitive architectures
  - Resurgence of neural networks (connectionism, backpropagation)
  - Valiant's PAC Learning Theory
  - Focus on experimental methodology
- 1990s
  - Data mining
  - Adaptive software agents and web applications
  - Text learning
  - Reinforcement learning (RL)
  - Inductive Logic Programming (ILP)
  - Ensembles: Bagging, Boosting, and Stacking
  - Bayes Net learning

History

- 2000s
  - Support vector machines
  - Kernel methods
  - Graphical models
  - Statistical relational learning
  - Transfer learning
  - Sequence labeling
  - Collective classification and structured outputs
  - Computer Systems Applications
    - Compilers
    - Debugging
    - Graphics
    - Security (intrusion, virus, and worm detection)
  - Email management
  - Personalized assistants that learn
  - Learning in robotics and vision

Definition of Machine Learning

- Machine learning/learning
  - A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P, if its performance at tasks in T, as measured by P, improves with experience E. [Tom Mitchell, Machine Learning]

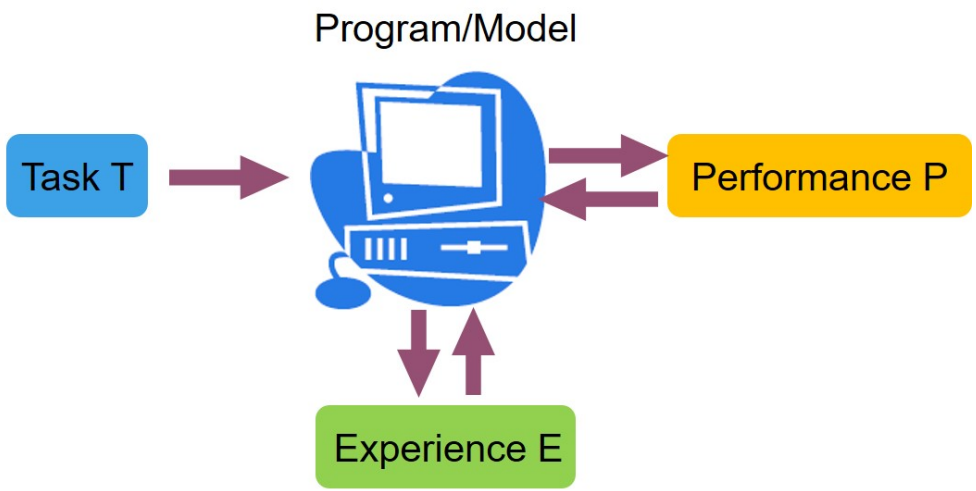

Improve on task T, with respect to performance metric P, based on experience E.

Example 1:

- T: Playing checkers
- P: Percentage of games won against an arbitrary opponent
- E: Playing practice games against itself

Example 2:

- T: Recognizing hand-written words
- P: Percentage of words correctly classified
- E: Database of human-labeled images of handwritten words

Example 3:

- T: Categorize email messages as spam or legitimate.
- P: Percentage of email messages correctly classified.
- E: Database of emails, some with human-given labels

Example 4:

- T: Driving on four-lane highways using vision sensors
- P: Average distance traveled before a human-judged error
- E: A sequence of images and steering commands recorded while observing a human driver.

Steps to Design a Learning System

- Choose the training experience
- Choose exactly what is to be learned, i.e. the target function.
- Choose how to represent the target function.
- Choose a learning algorithm to infer the target function from the experience.

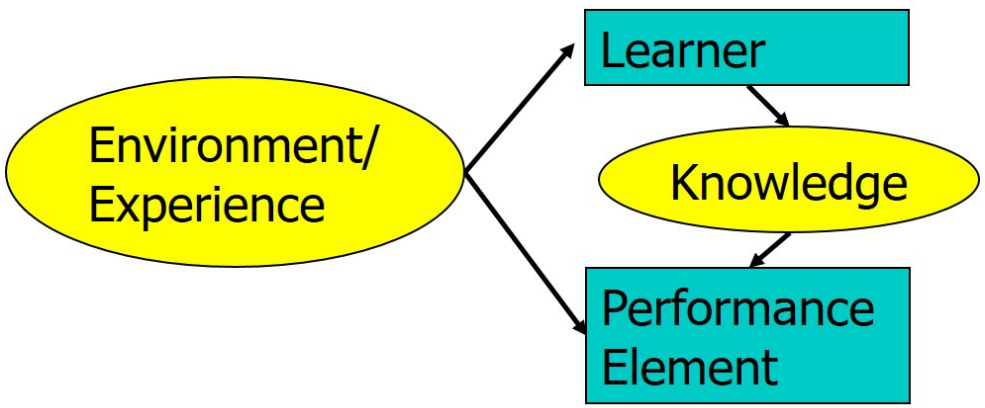

Training Experience

- Direct experience: Given sample input and output pairs for a useful target function.
  - Checker boards labeled with the correct move, e.g. extracted from record of expert play
- Indirect experience: Given feedback which is not direct I/O pairs for a useful target function.
  - Potentially arbitrary sequences of game moves and their final game results.
- Credit/Blame Assignment Problem: How to assign credit blame to individual moves given only indirect feedback?

Source of Training Data

- Provided random examples outside of the learner's control.
  - Negative examples available or only positive?
- Good training examples selected by a "benevolent" teacher.
  - "Near miss" examples
- Learner can query an oracle about class of an unlabeled example in the environment.
- Learner can construct an arbitrary example and query an oracle for its label.
- Learner can design and run experiments directly in the environment without any human guidance.

Applications of Machine Learning

Document Search

- Given counts of words in a document, determine what its topic is.
- Group documents by topic without a pre-specified list of topics.
- Many words in a document, many, many documents available on the web.

Image/Video Understanding

- Given an/a image/video, determine what objects it contains.
- Determine what semantics it contains
- Determine what actions it contains.

Cancer Diagnosis

- Given data on expression levels of genes, classify the type of tumor.
- Discover categories of tumors having different characteristics.

Marketing

- Given data on age, income, etc., predict how much each customers spends.
- Discover how the spending behaviors of customers are related.
- Fair amount of data on each customer, but messy
- May have data on a very large number of customer.

Example 1: Handwritten Digit Recognition

- Handcrafted rules will result in large number of rules and exceptions
- Better to have a machine that learns from a large training set
- Handwriting recognition cannot be done without machine learning!

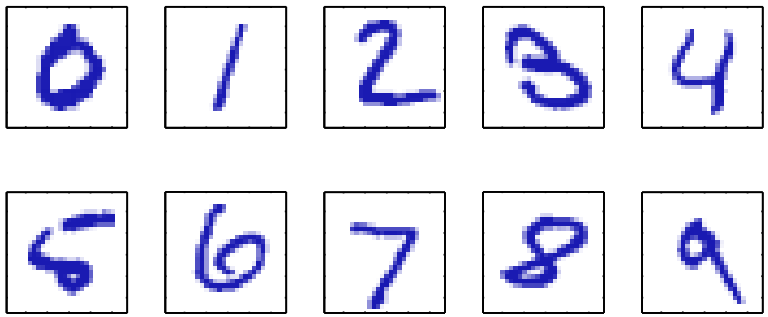

Everyone has different writing style!

Example 2: Autonomous Driving-ALVINN

Drives 70 mph on a public highway

Predecessor of the Google car

Camera image

- 30 outputs for steering
- 4 hidden units
- 30x32 pixels as inputs

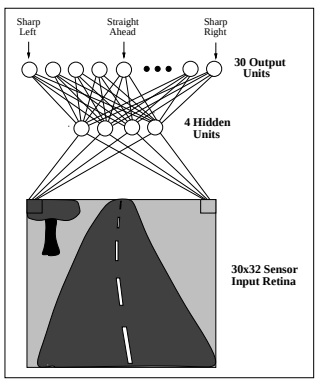

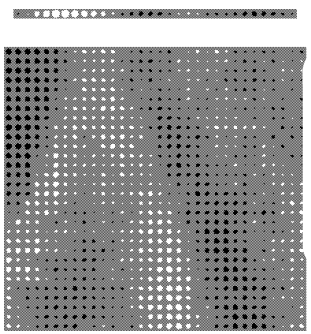

30x32 weights into one out of four hidden unit

Example 3: Breast Cancer Diagnosis

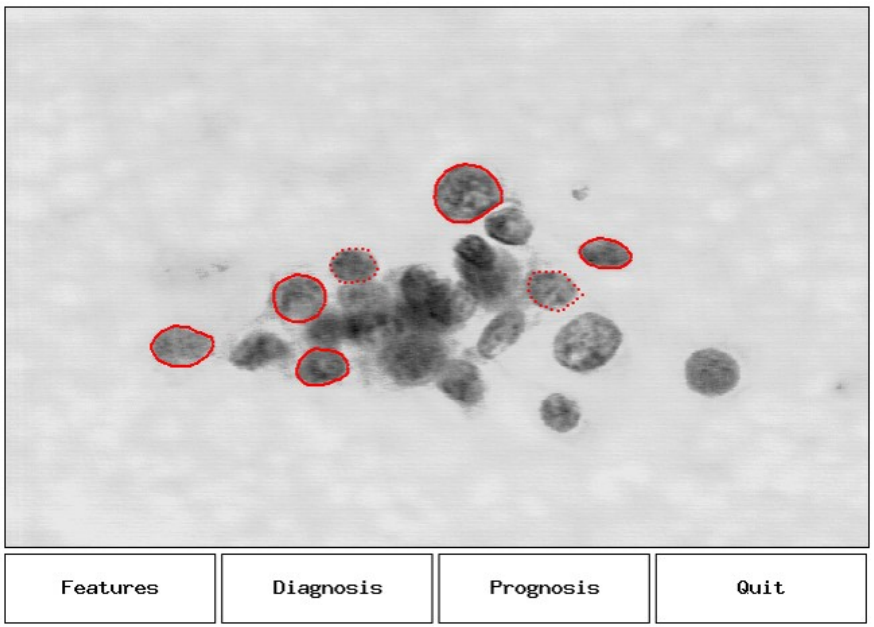

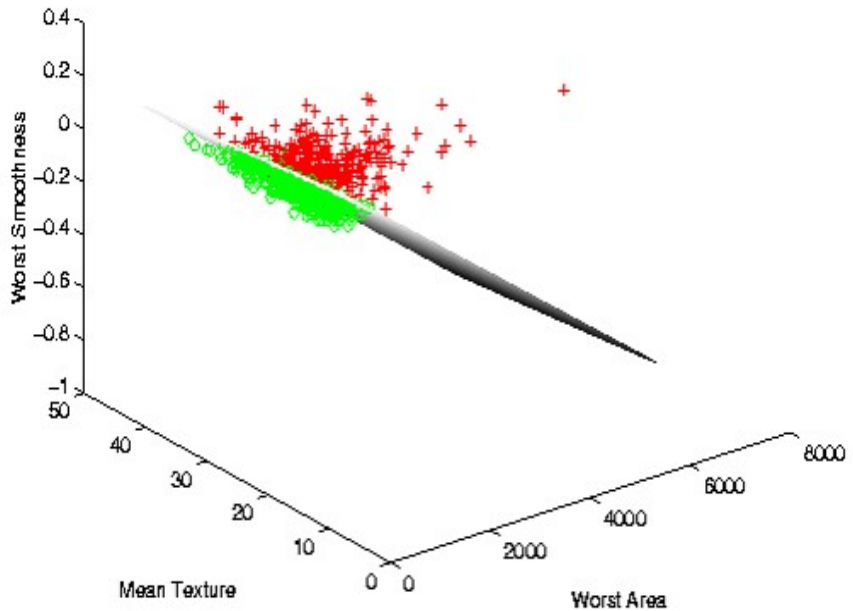

Research by Mangasarian,Street, Wolberg

Why is machine learning necessary?

1. Learning is a hallmark of intelligence; many would argue that a system that cannot learn is not intelligent.
2. Without learning, everything is new; a system that cannot learn is not efficient because it rederives each solution and repeatedly makes the same mistakes.

Why is learning possible?

Because there are regularities in the world.

Categories of Machine Learning

1. Supervised learning: learning with a teacher
   I.e., training examples with labels are given.
2. Unsupervised learning: learning without a teacher
   I.e., training examples without labels.
3. Reinforcement Learning: learning by interacting
4. Semi-supervised learning: partially supervised learning
5. Active learning: actively making queries

Supervised Learning

- In the ML literature, a supervised learning problem has the following characteristics:
  - We are primarily interested in prediction.
  - We are interested in predicting only one thing.
  - The possible values of what we want to predict are specified, and we have some training cases for which its value is known.
- The thing we want to predict is called the target or the response variable.
- Usually, we need training data

Supervised Learning

- For classification problem, we want to predict the class of an item.
  - The type of tumor, the topic of a document, the semantics contained in an image, whether a customer will purchase a product.
- For a regression problem, we want to predict a numerical quantity.
  - The amount of customer spends, the blood pressure of a patient, etc.
- To make predictions, we have various inputs,
  - Gene expression levels for predicting tumor type, age and income for predicting amount spent, the features of images with known semantics

Classification and Regression

- Classification: finding decision boundaries

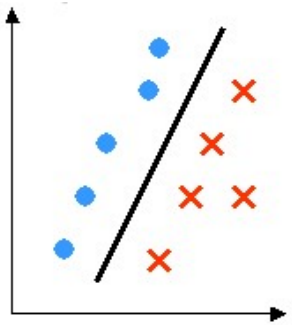

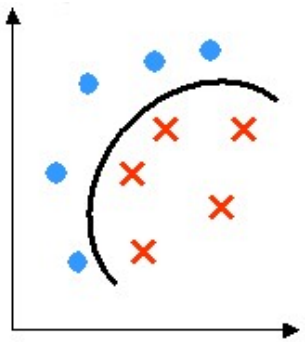

- Regression: fitting a curve/plane to data

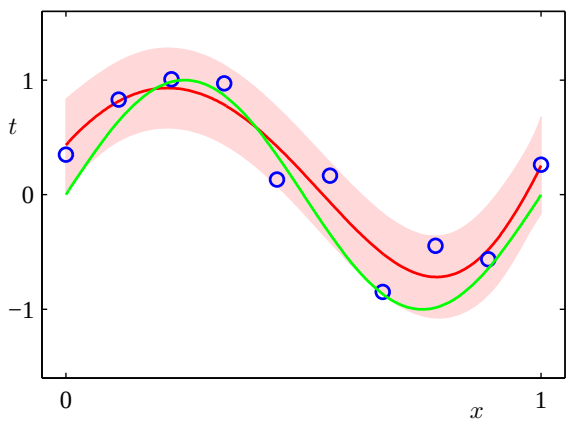

Supervised Classification Problem

- Cancer diagnosis (Training Set)

Patient ID	# of Tumors	Avg Area	Avg Density	Diagnosis
1	5	20	118	Malignant
2	3	15	130	Benign
3	7	10	52	Benign
4	2	30	100	Malignant

- Use the above training set to learn how to classify patients where diagnosis is not known (Test Set):

Patient ID	# of Tumors	Avg Area	Avg Density	Diagnosis
101	4	16	95	?
102	9	22	125	?
103	1	14	80	?

- The input data is often easily obtained, whereas the classification is not.

How to Make Predictions?

- Main methods
  - We can train a model by using the training data to estimate parameters of it
  - Use these parameters to make predictions for the test data.
  - Such approaches save computation when we make predictions for test data.
  - That is, estimate parameters once, use them many times.
  - e.g. Linear regression

$$\hat{y} = \hat{\beta}_0 + \sum_{j=1}^{p} \hat{\beta}_j x_j$$

- Other methods
  - Nearest-Neighbor like methods
  - Need to store training data

Nearest-Neighbor Methods

- Make predictions for test data based on a subset of training cases, e.g., by approximating the mean, median or mode of $P(y|x)$.

$$\hat{y} = \frac{1}{K} \sum_{i \in N_K(x)} y_i$$

- Big question: How to choose K?
- If K is too small, we may "overfitting", but if K is too big, we will average over training cases that aren't relevant to the test case.

Comparisons

- These two methods are opposite w.r.t. computation.
  - NN like methods are memory-based methods. We need to remember all the training data.
  - Linear regression, after getting parameters, can forget the training data, and just use the parameters.
- They are also opposite w.r.t. to statistical properties.
  - NN makes few assumptions about the data, and has a high potential for overfitting.
  - Linear regression makes strong assumption about the data, and consequently has a high potential for bias.

Unsupervised Learning

- For an unsupervised learning problem, we do not focus on prediction of any particular thing, but rather try to find interesting aspects of the data.
- Examples:
  - We may find clusters of patients with similar symptoms, which we call diseases.
  - We may find clusters of large number of images.

Unsupervised Learning: Clustering

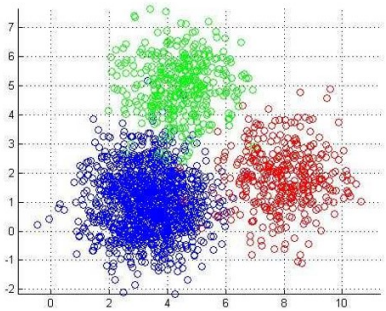

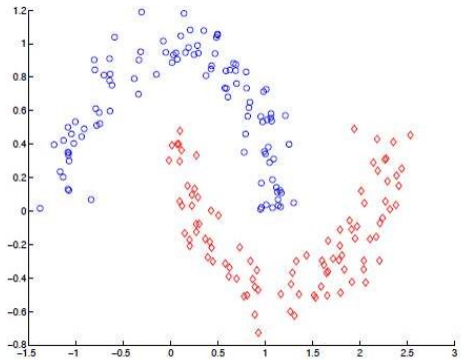

Reinforcement Learning

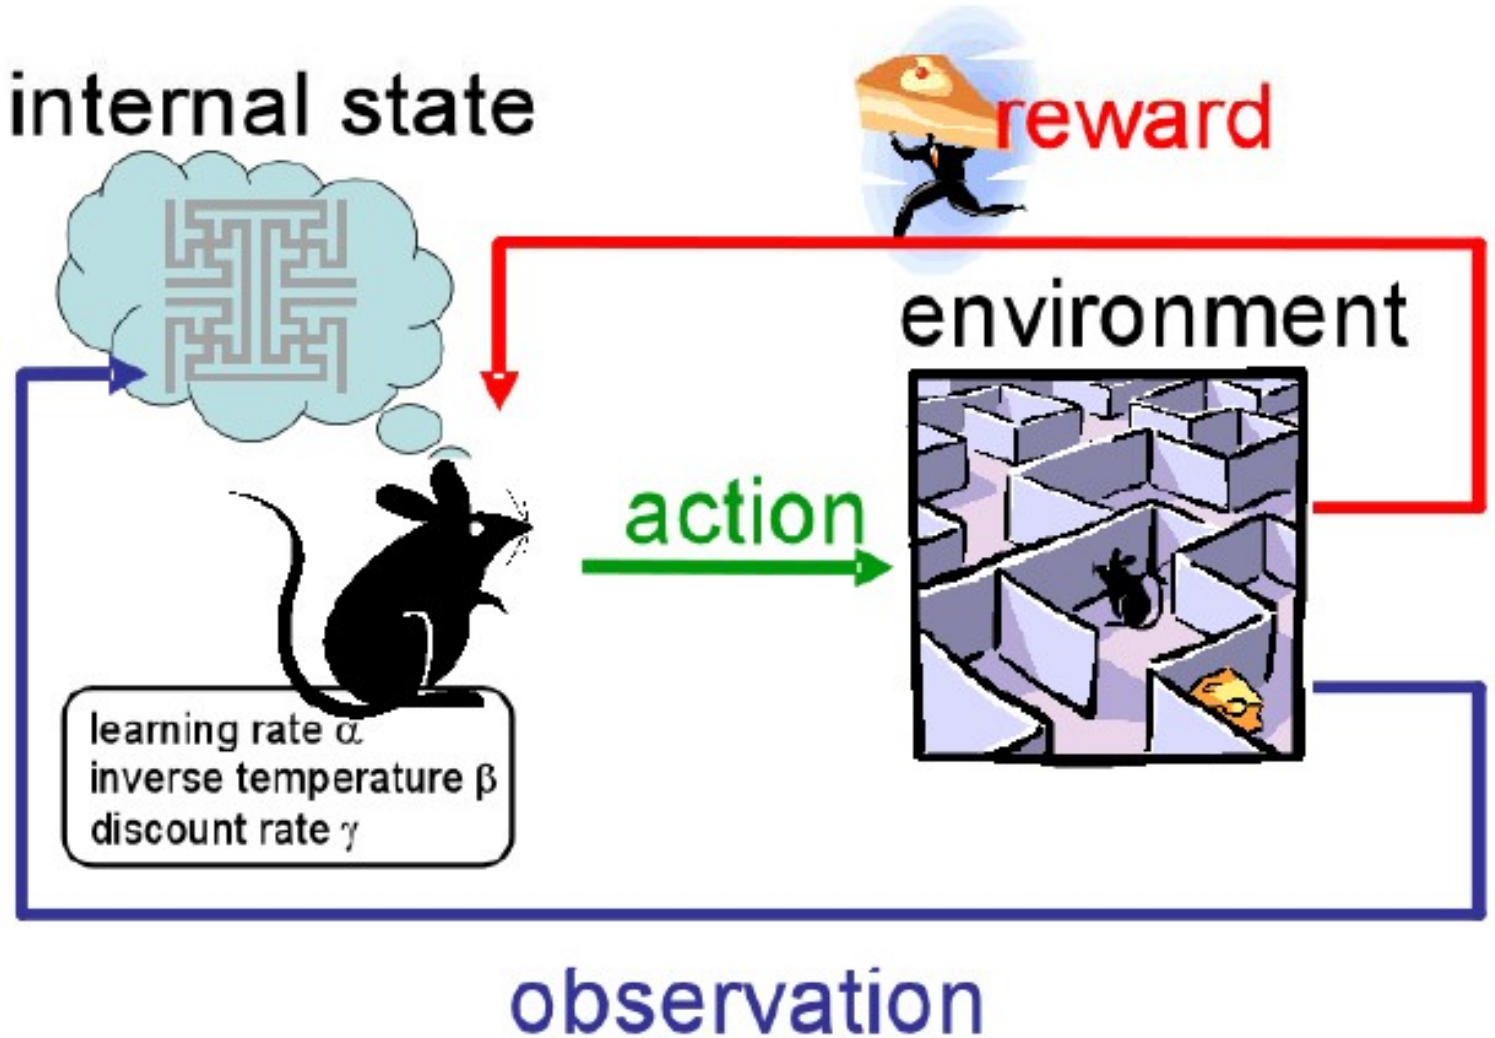

Reinforcement Learning

- Learning from interaction (with environment)
- Goal-directed learning
- Learning what to do and its effect
- Trial-and-error search and delayed reward
  The two most important distinguishing features of reinforcement learning
- The agent has to exploit what it already knows in order to obtain reward, but it also has to explore in order to make better action selections in the future.
- Dilemma: neither exploitation nor exploration can be pursued exclusively without failing at the task.

Reinforcement Learning

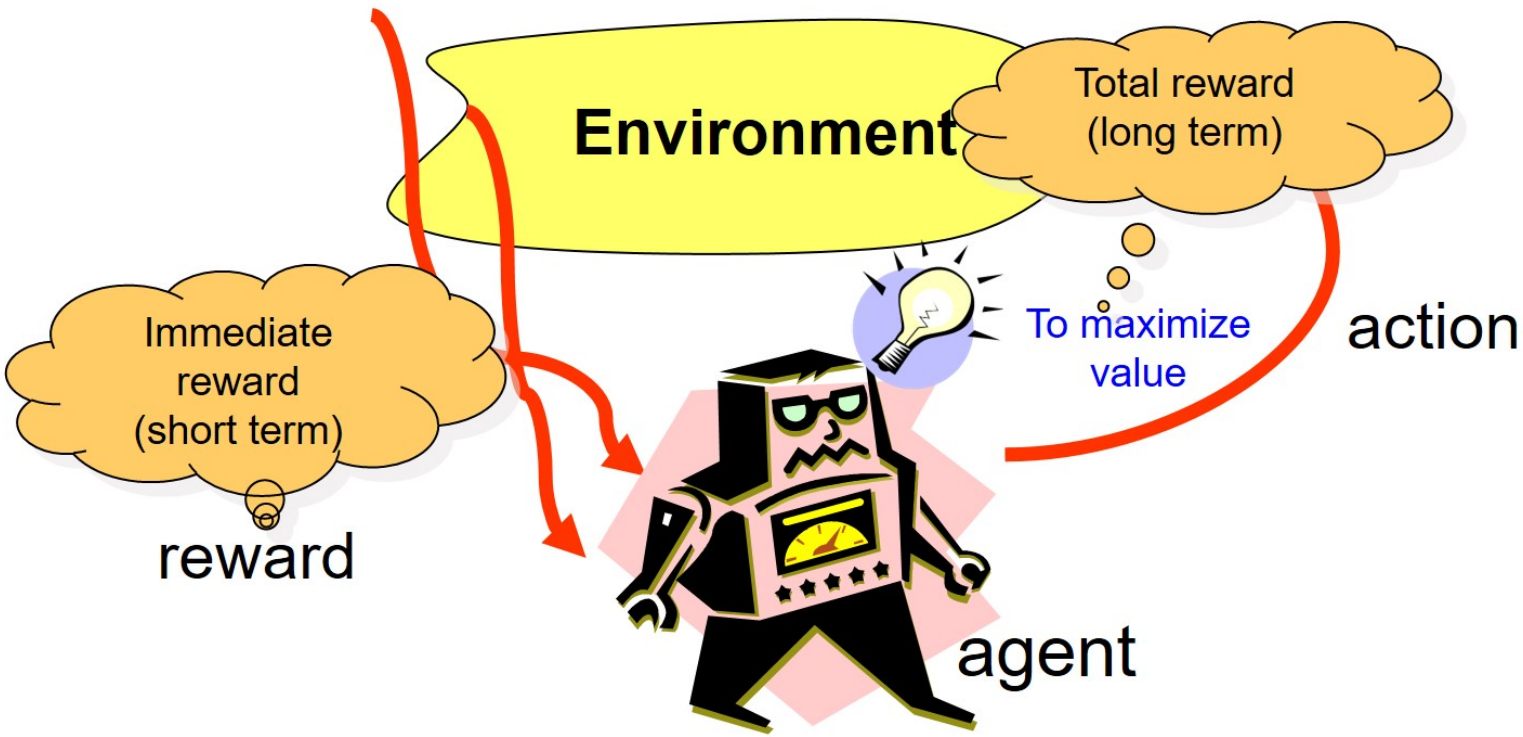

Reinforcement Learning

- Example (Bioreactor)
- State
  - current temperature and other sensory readings, composition, target chemical
- Actions
  - how much heating, stirring are required?
  - what ingredients need to be added?
- Reward
  - moment-by-moment production of desired chemical

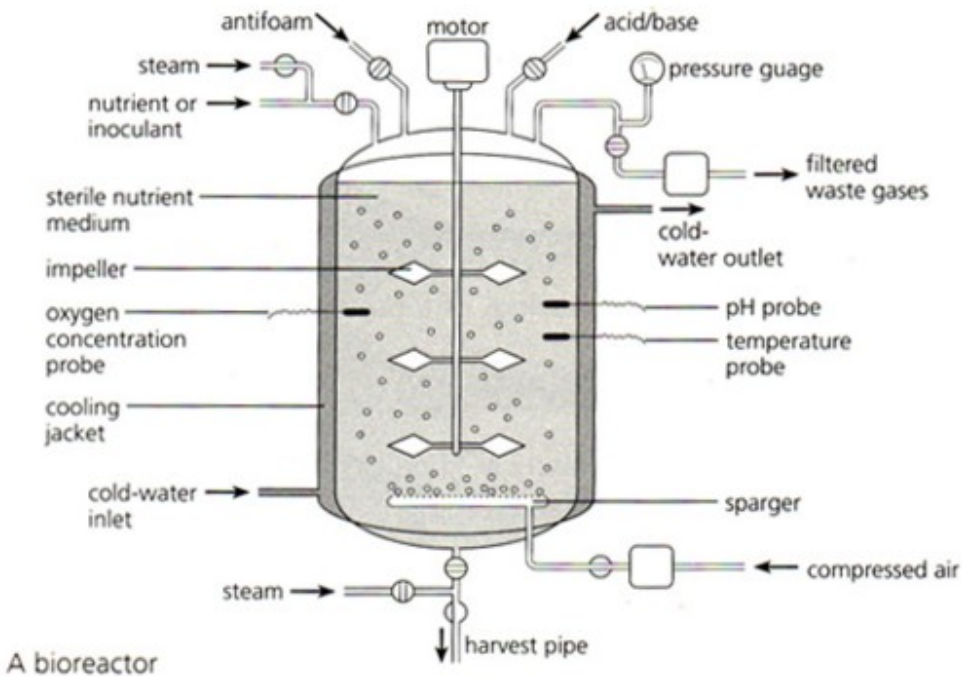

A bioreactor

Reinforcement Learning

- Example (Pick-and-Place Robot)
- State
  - current positions and velocities of joints
- Actions
  - voltages to apply to motors
- Reward
  - reach end-position successfully, speed, smoothness of trajectory

Reinforcement Learning

- Example (Recycling Robot)
- State
  - charge level of battery
- Actions
  - look for cans, wait for can, go recharge
- Reward
  - positive for finding cans, negative for running out of battery

Semi-supervised Learning

- As the name suggests, it is in between Supervised and Unsupervised learning techniques w.r.t the amount of labelled and unlabelled data required for training.
- With the goal of reducing the amount of supervision required compared to supervised learning.
- At the same time, improving the results of unsupervised clustering to the expectations of the user.

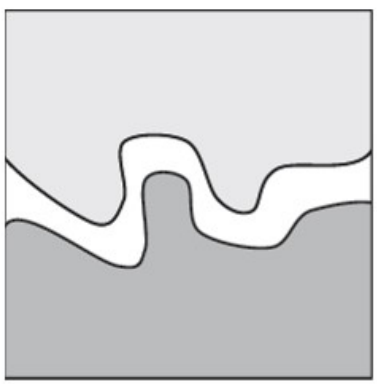
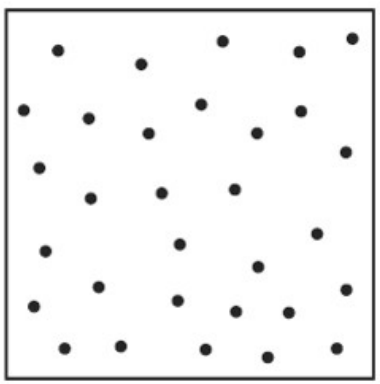
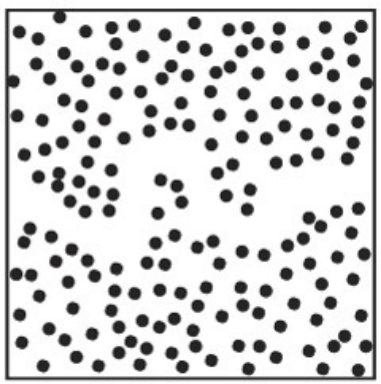

With lots of unlabeled data the decision boundary becomes apparent.

Overview of Semi-Supervised Learning

- Constrained Clustering
- Distance Metric Learning
- Manifold based Learning
- Sparsity based Learning (Compressed Sensing).
- Active Learning

Constrained Clustering

- When we have any of the following:
  - Class labels for a subset of the data.
  - Domain knowledge about the clusters.
  - Information about the 'similarity' between objects.
  - User preferences.
- May be pairwise constraints or a labeled subset.
  - Must-link or cannot-link constraints.
  - Labels can always be converted to pairwise relations.
- Can be clustered by searching for partitionings that respect the constraints.
- Recently the trend is toward similarity-based approaches.

Constrained Clustering

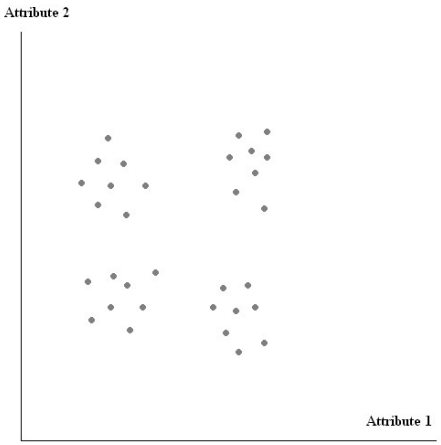

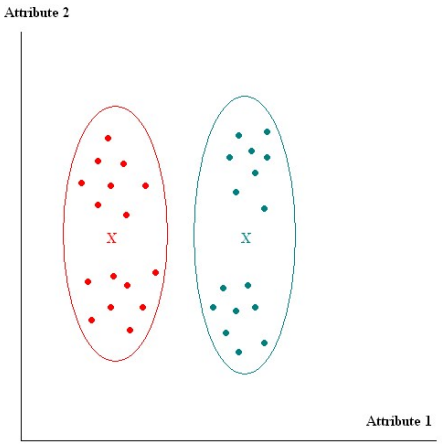

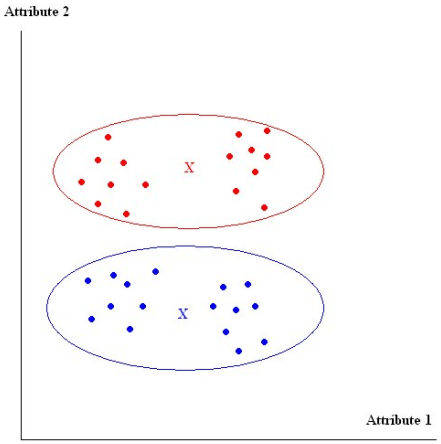

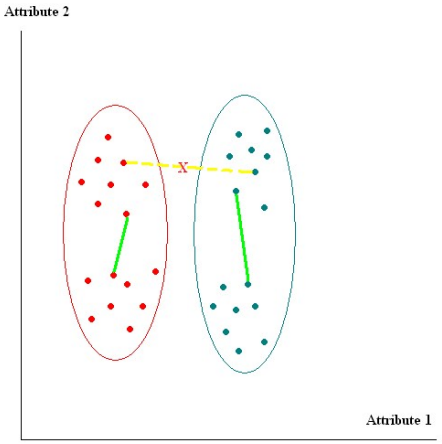

Active Learning

- Basic idea:
  - Traditional supervised learning algorithms passively accept training data.
  - Instead, query for annotations on informative images from the unlabeled data.
  - Theoretical results show that large reductions in training sizes can be obtained with active learning!
- But how to find images that are the most informative?

Active Learning

- One idea uses uncertainty sampling.
- Images on which you are uncertain about classification might be informative!
- What is the notion of uncertainty?
  - Idea: Train a classifier like SVM on the training set.
  - For each unlabeled image, output probabilities indicating class membership.
  - Estimate probabilities can be used to infer uncertainty.
  - A one-vs-one SVM approach can be used to tackle multiple classes.

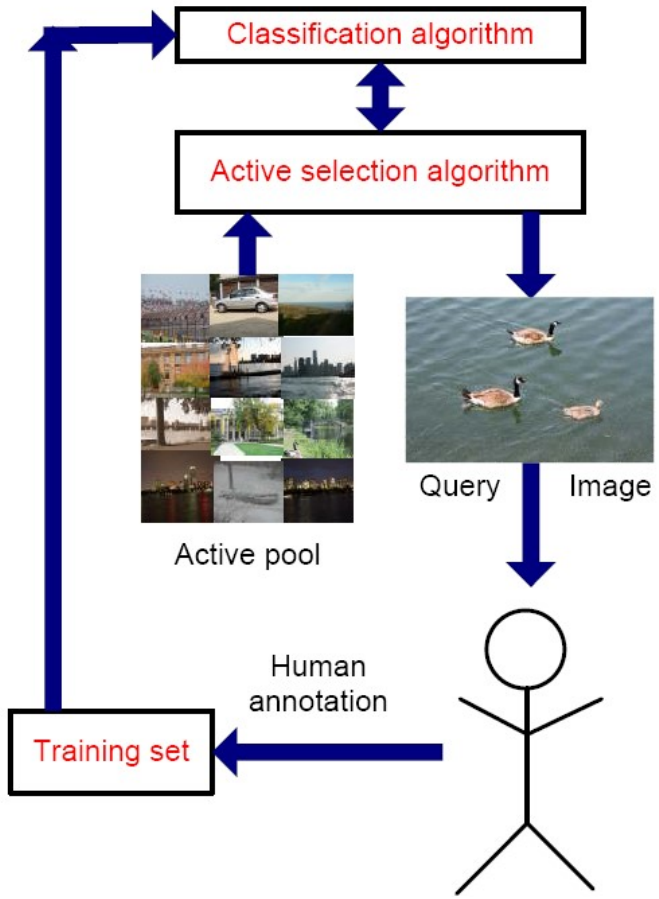

Challenges for ML

- Handling complexity
  - Involve many variables, how can we handle this complexity without getting into trouble.
- Optimization and Integration
  - Usually involve finding the best values for some parameters (an optimization problem), or average over many plausible values (an integration problem). How can we do this efficiently when there are many parameters.
- Visualization
  - Understanding what's happening is hard, 2D? 3D?
- All these challenges become greater when there are many variables or parameters —-the so-called "curse of dimensionality".
  - But more variables also provide more information
  - A blessing? A curse?

How to handle complexity

- Properly dealing with complexity is a crucial issue for machine learning.
- Limiting complexity is one approach
  - Use a model that is complex enough to represent the essential aspects of the problem, but that is not so complex that overfitting occurs.
  - Overfitting happens when we choose parameters of a model that fit the data we have very well, but do poorly on new data (poor generalization ability).
  - Cross-validation, regularization,
- Reducing dimensionality is another possibility.
  - It is apparent that things become simpler if can find out how to reduce the large number of variables to a small number.
- Averaging over complexity is the Bayesian approach.
  - Use as complex a model might be needed, but don't choose a single parameter values. Instead, average the predictions found using all the parameter values that fit the data reasonably well, and which are plausible for the problem

Example of Complexity

Plots of polynomials having various degree, shown as red curves.

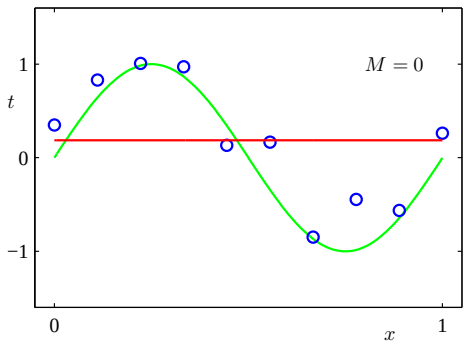

Figure: Degree = 0

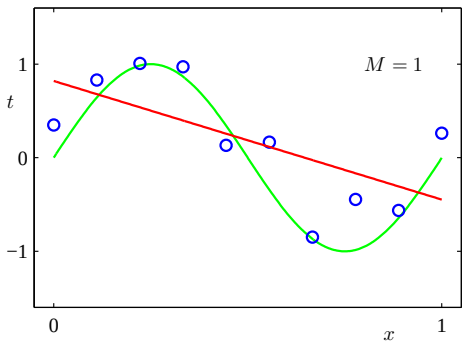

Figure: Degree = 1

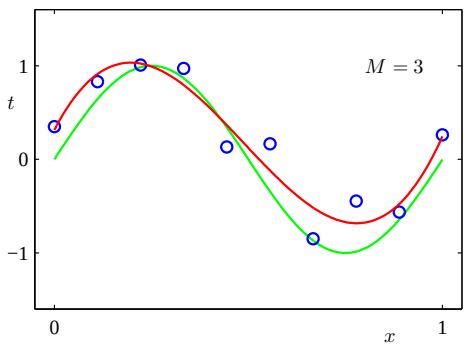

Figure: Degree = 3

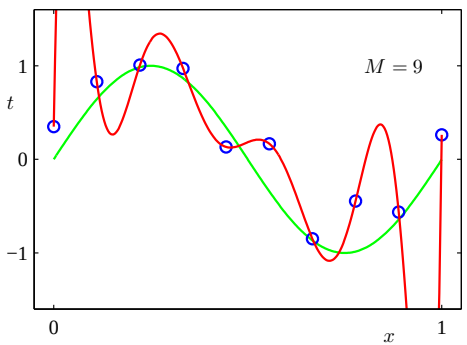

Figure: Degree = 9

Example of Complexity

Graphs of the root-square error, evaluated on the training set and on an independent test set for various degree.

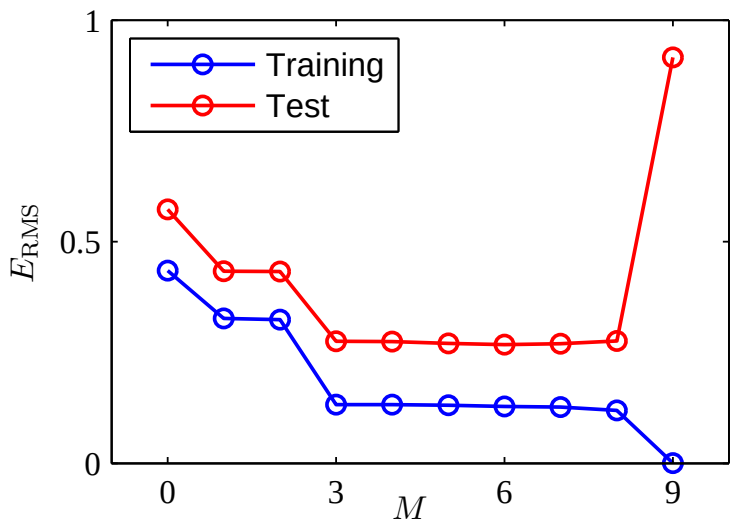

Does Complexity should be limited?

- If we make predictions using "the best" parameters of a model, we have to limit the number of parameters to avoid overfitting.
- For this example, the model with degree=3 seems good. We might be able to choose a good value for M using the method of "cross validation", which looks for the value that does best at prediction one part of the data from the rest of the data.
- But we know $\sin(2\pi x)$ is not a polynomial function, it has an infinite series representation with terms of arbitrarily high order.
- How can it be good to use a model that we know is false?
  - The Bayesian answer: It is not good. We should abandon the idea of using the best parameters and instead average over all plausible values for the parameters. Then we can use a model (perhaps a very complex one) that is as close to being correct as we can manage.

Reducing Dimensionality

- Suppose dimension of input data is 1000, can we replace these with fewer ones, without loss of information.
- On simple way is to use PCA (Principal Component Analysis)
  - Suppose that all data are in a space, we first find the direction of highest variance of these data points, then the direction of second-highest variance that is orthogonal to the first one, so on and so forth
  - Replace each training sample by the projections of the inputs on some directions.
- Might discard useful info., but keep most of it.

Thanks!

Q & A