

# Machine Learning

## Lecture 8: Factor Analysis and Principle Component Analysis

**Feng Li**

`fli@sdu.edu.cn`

`https://funglee.github.io`

School of Computer Science and Technology  
Shandong University

Fall 2018

# Higher Dimension But Less Data

- Consider a case with  $n \gg m$ 
  - The given training data span only a low-dimensional subspace of  $\mathbb{R}^n$
- Model the data as Gaussian and estimate the mean and covariance using MLE

$$\mu = \frac{1}{m} \sum_{i=1}^m x^{(i)}$$

$$\Sigma = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu)(x^{(i)} - \mu)^T$$

- $\Sigma$  may be singular such that  $\Sigma^{-1}$  does not exist and  $1/|\Sigma|^{1/2} = 1/0$

$$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp \left( -\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right)$$

# Marginals and Conditionals of Gaussians

- Consider a vector-valued random variable

$$x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

where  $x_1 \in \mathbb{R}^r$ ,  $x_2 \in \mathbb{R}^s$  and  $x \in \mathbb{R}^{r+s}$

- $x$  follows a Gaussian distribution  $x \sim \mathcal{N}(\mu, \Sigma)$

$$\mu = \begin{bmatrix} \mu_1 \\ \mu_2 \end{bmatrix} \quad \Sigma = \begin{bmatrix} \Sigma_{11} & \Sigma_{12} \\ \Sigma_{21} & \Sigma_{22} \end{bmatrix}$$

where  $\mu_1 \in \mathbb{R}^r$ ,  $\mu_2 \in \mathbb{R}^s$ ,  $\Sigma_{11} \in \mathbb{R}^{r \times r}$ ,  $\Sigma_{12} \in \mathbb{R}^{r \times s}$ ,  $\Sigma_{21} \in \mathbb{R}^{s \times r}$ , and  $\Sigma_{22} \in \mathbb{R}^{s \times s}$

- Also,  $\Sigma_{12} = \Sigma_{21}^T$  due to the symmetry of  $\Sigma$

# Marginals and Conditionals of Gaussians (Contd.)

- $x_1$  and  $x_2$  are jointly multivariate Gaussian
- What is the marginal distribution of  $x_1$ 
  - $E[x_1] = \mu_1$
  - $Cov(x_1) = \Sigma_{11}$
- Since the marginal distribution of Gaussian are themselves Gaussian, we have

$$x_1 \sim \mathcal{N}(\mu_1, \Sigma_{11})$$

- Conditional multivariate Gaussian distribution  $x_1 \mid x_2 \sim \mathcal{N}(\mu_{1|2}, \Sigma_{1|2})$

$$\mu_{1|2} = \mu_1 + \Sigma_{12}\Sigma_{22}^{-1}(x_2 - \mu_2)$$

$$\Sigma_{1|2} = \Sigma_{11} - \Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21}$$

# Factor Analysis Model

- $x = \mu + \Lambda z + \varepsilon$ 
  - $x \in \mathbb{R}^n$ ,  $\mu \in \mathbb{R}^n$ ,  $\Lambda \in \mathbb{R}^{n \times k}$ ,  $z \in \mathbb{R}^k$ ,  $\varepsilon \in \mathcal{R}^n$
  - $\Lambda$  is the factor loading matrix
  - $z \sim \mathcal{N}(0, I)$  (zero-mean independent normals, with unit variance)
  - $\varepsilon \sim \mathcal{N}(0, \Psi)$  where  $\Psi$  is a diagonal matrix (the observed variables are independent given the factors)
- How do we get the training data  $\{x^{(i)}\}_i$ ?
  - Generate  $\{z^{(i)}\}_i$  according to a multivariate Gaussian distribution  $\mathcal{N}(0, I)$
  - Map  $\{z^{(i)}\}_i$  into a  $n$ -dimensional affine space by  $\Lambda$  and  $\mu$
  - Generate  $\{x^{(i)}\}_i$  by sampling the above affine space with noise  $\varepsilon$
- Equivalently,

$$z \sim \mathcal{N}(0, I)$$

$$x|z \sim \mathcal{N}(\mu + \Lambda z, \Psi)$$

## Factor Analysis Model (Contd.)

- $z$  and  $x$  have a joint Gaussian distribution

$$\begin{bmatrix} z \\ x \end{bmatrix} \sim \mathcal{N}(\mu_{zx}, \Sigma)$$

- Question: How to calculate  $\mu_{zx}$  and  $\Sigma$ ?
- Since  $E[z] = 0$ , we have

$$E[x] = E[\mu + \Lambda z + \epsilon] = \mu + \Lambda E[z] + E[\epsilon] = \mu$$

and then

$$\mu_{zx} = \begin{bmatrix} \vec{0} \\ \mu \end{bmatrix}$$

## Factor Analysis Model (Contd.)

- Since  $z \sim \mathcal{N}(0, I)$ , we have

$$\Sigma_{zz} = E[(z - E[z])(z - E[z])^T] = \text{Cov}(z) = I$$

$$\Sigma_{zx} = E[(z - E[z])(x - E[x])^T] = E[z(\mu + \Lambda z + \epsilon - \mu)^T] = \Lambda^T$$

$$\Sigma_{xx} = E[(\mu + \Lambda z + \epsilon - \mu)(\mu + \Lambda z + \epsilon - \mu)^T] = \Lambda\Lambda^T + \Psi$$

- Putting everything together, we therefore have

$$\begin{bmatrix} z \\ x \end{bmatrix} \sim \mathcal{N} \left( \begin{bmatrix} \vec{0} \\ \mu \end{bmatrix}, \begin{bmatrix} I & \Lambda^T \\ \Lambda & \Lambda\Lambda^T + \Psi \end{bmatrix} \right)$$

- Then,  $x \sim \mathcal{N}(\mu, \Lambda\Lambda^T + \Psi)$
- Log-likelihood function

$$\ell(\mu, \Lambda, \Psi) = \log \prod_{i=1}^m \frac{1}{(2\pi)^{n/2} |\Sigma_{xx}|^{1/2}} \exp \left( -\frac{1}{2} (x^{(i)} - \mu)^T \Sigma_{xx}^{-1} (x^{(i)} - \mu) \right)$$

# EM Algorithm Review

- Repeat the following step until convergence
  - (E-step) For each  $i$ , set

$$Q_i(z^{(i)}) := p(z^{(i)} \mid x^{(i)}; \theta)$$

- (M-step) set

$$\theta := \arg \max_{\theta} \sum_i \sum_{z^{(i)}} Q_i(z^{(i)}) \log \frac{p(x^{(i)}, z^{(i)}; \theta)}{Q_i(z^{(i)})}$$



# EM Algorithm for Factor Analysis

- Recall that if

$$\begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \sim \mathcal{N}\left(\mu = \begin{bmatrix} \mu_1 \\ \mu_2 \end{bmatrix}, \Sigma = \begin{bmatrix} \Sigma_{11} & \Sigma_{12} \\ \Sigma_{21} & \Sigma_{22} \end{bmatrix}\right)$$

we then have

$$x_1|x_2 \sim \mathcal{N}(\mu_{1|2}, \Sigma_{1|2})$$

where

$$\mu_{1|2} = \mu_1 + \Sigma_{12}\Sigma_{22}^{-1}(x_2 - \mu_2)$$

$$\Sigma_{1|2} = \Sigma_{11} - \Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21}$$

## EM Algorithm for Factor Analysis (Contd.)

- Since

$$\begin{bmatrix} z \\ x \end{bmatrix} \sim \mathcal{N} \left( \begin{bmatrix} \vec{0} \\ \mu \end{bmatrix}, \begin{bmatrix} I & \Lambda^T \\ \Lambda & \Lambda\Lambda^T + \Psi \end{bmatrix} \right)$$

we have

$$z^{(i)} | x^{(i)}; \mu, \Lambda, \Psi \sim \mathcal{N}(\mu_{z^{(i)} | x^{(i)}}, \Sigma_{z^{(i)} | x^{(i)}})$$

where

$$\mu_{z^{(i)} | x^{(i)}} = \Lambda^T (\Lambda\Lambda^T + \Psi)^{-1} (x^{(i)} - \mu)$$

$$\Sigma_{z^{(i)} | x^{(i)}} = I - \Lambda^T (\Lambda\Lambda^T + \Psi)^{-1} \Lambda$$

- Calculate  $Q_i(z^{(i)})$  in the E-step

$$Q_i(z^{(i)}) = \frac{\exp \left( -\frac{1}{2} (z^{(i)} - \mu_{z^{(i)} | x^{(i)}})^T \Sigma_{z^{(i)} | x^{(i)}}^{-1} (z^{(i)} - \mu_{z^{(i)} | x^{(i)}}) \right)}{(2\pi)^{k/2} |\Sigma_{z^{(i)} | x^{(i)}}|^{1/2}}$$

## EM Algorithm for Factor Analysis (Contd.)

- In M-step, we maximize

$$\sum_{i=1}^m \int_{z^{(i)}} Q_i(z^{(i)}) \log \frac{p(x^{(i)}, z^{(i)}; \mu, \Lambda, \Psi)}{Q_i(z^{(i)})} dz^{(i)}$$

with respect to  $\mu$ ,  $\Lambda$ , and  $\Psi$

- Results are as follows



$$\mu = \frac{1}{m} \sum_{i=1}^m x^{(i)}$$

$$\Lambda = \left( \sum_{i=1}^m (x^{(i)} - \mu) \mu_{z^{(i)}|x^{(i)}}^T \right) \left( \sum_{i=1}^m \mu_{z^{(i)}|x^{(i)}} \mu_{z^{(i)}|x^{(i)}}^T + \Sigma_{z^{(i)}|x^{(i)}} \right)^{-1}$$

$$\Phi = \text{diag} \left( \frac{1}{m} \sum_{i=1}^m x^{(i)} x^{(i)T} - x^{(i)} \mu_{z^{(i)}|x^{(i)}}^T \Lambda^T - \Lambda \mu_{z^{(i)}|x^{(i)}} x^{(i)T} + \right. \\ \left. \Lambda (\mu_{z^{(i)}|x^{(i)}} \mu_{z^{(i)}|x^{(i)}}^T + \Sigma_{z^{(i)}|x^{(i)}}) \Lambda^T \right)$$

# Dimensionality Reduction


- Usually considered an unsupervised learning method
- Used for learning the low-dimensional structures in the data



- Also useful for “feature learning” or “representation learning” (learning a better, often smaller-dimensional, representation of the data), e.g.,
  - Documents using using topic vectors instead of bag-of-words vectors
  - Images using their constituent parts (faces - eigenfaces)
- Can be used for speeding up learning algorithms

# Dimensionality Reduction

- Exponentially large # of examples required to “fill up” high-dim spaces



- Fewer dimensions  $\Rightarrow$  Less chances of overfitting  $\Rightarrow$  Better generalization
- Dimensionality reduction is a way to beat the curse of dimensionality

# Linear Dimensionality Reduction

- A projection matrix  $U = [u_1 u_2 \cdots u_K]$  of size  $D \times K$  defines  $K$  linear projection direction
- Use  $U$  to transform  $x^{(i)} \in \mathbb{R}^D$  into  $z^{(i)} \in \mathbb{R}^K$

$$\begin{matrix} K \times 1 \\ \text{orange bar} \\ z_i \end{matrix} = \begin{matrix} K \times D \\ \text{gray box} \\ \begin{matrix} \text{--- } u_1^T \text{ ---} \\ U \\ \text{--- } u_K^T \text{ ---} \end{matrix} \end{matrix} * \begin{matrix} \text{blue bar} \\ x^{(i)} \end{matrix}$$

- $z^{(i)} = U^T x^{(i)} = [u_1^T x^{(i)}, u_2^T x^{(i)}, \cdots u_K^T x^{(i)}]^T$  is a  $K$ -dim projection of  $x^{(i)}$ 
  - $z^{(i)} \in \mathbb{R}^K$  is also called low-dimensional "embedding" of  $x^{(i)} \in \mathbb{R}^D$

# Linear Dimensionality Reduction


- $X = [x^{(1)} \ x^{(2)} \ \dots \ x^{(N)}]$  is  $D \times N$  matrix denoting all the  $N$  data points
- $Z = [z^{(1)} \ z^{(2)} \ \dots \ z^{(N)}]$  is  $K \times N$  matrix denoting embeddings of the data points
- With this notation, the figure on previous slide can be re-drawn as

$$\begin{array}{c} K \times N \\ \boxed{\mathbf{Z}} \end{array} = \begin{array}{c} K \times D \\ \boxed{\mathbf{U}^T} \end{array} * \begin{array}{c} D \times N \\ \boxed{\mathbf{X}} \end{array}$$

- How do we learn the “best” projection matrix  $U$ ?
- What criteria should we optimize for when learning  $U$
- Principle Component Analysis (PCA) is an algorithm for doing this

# Principle Component Analysis (PCA)


- PCA is a technique widely used for applications such as dimensionality reduction, lossy data compression, feature extraction, and data visualization
- Two commonly used definitions
  - Learning projection directions that capture maximum variance in data
  - Learning projection directions that result in smallest reconstruction error
- Can also be seen as changing the basis in which the data is represented (and transforming the features such that new features become decorrelated)





## Variance Captured by Projections

- Consider  $x^{(i)} \in \mathbb{R}^D$  on a one-dim subspace defined by  $u_1 \in \mathbb{R}^D$
- Projection of  $x^{(i)}$  along a one-dim subspace  $u_1 = u_1^T x^{(i)}$



- Mean of projections of all the data

$$\frac{1}{N} \sum_{i=1}^N u_1^T x^{(i)} = u_1^T \frac{1}{N} \sum_{i=1}^N x^{(i)} = u_1^T \mu$$

# Variance Captured by Projections


- Variance of the projected data

$$\frac{1}{N} \sum_{i=1}^N (u_1^T x^{(i)} - u_1^T \mu)^2 = \frac{1}{N} \sum_{i=1}^N [u_1^T (x^{(i)} - \mu)]^2 = u_1^T S u_1$$

- $S$  is the  $D \times D$  data covariance matrix

$$S = \frac{1}{N} \sum_{i=1}^N (x^{(i)} - \mu)(x^{(i)} - \mu)^T$$

- Variance of the projected data ("spread" of the yellow points)
- If data already centered at  $\mu = 0$ , then  $S = \frac{1}{N} \sum_{i=1}^N x^{(i)}(x^{(i)})^T$



# Optimization Problem

- We want  $u_1$  s.t. the variance of the projected data is maximized

$$\max_{u_1} u_1^T S u_1$$

- To prevent trivial solution (max var. = infinite), assume  $\|u_1\| = 1 = u_1^T u_1$
- The method of Lagrange multipliers

$$\mathcal{L}(u_1, \lambda_1) = u_1^T S u_1 + \lambda_1(1 - u_1^T u_1)$$

where  $\lambda_1$  is a Lagrange multiplier

- If  $u_1^*$  is the optimal solution for the original constrained problem, then there exists  $\lambda_1^*$  such that  $(u_1^*, \lambda_1^*)$  is a stationary point for the Lagrange function (stationary points are those points where the partial derivatives of  $\mathcal{L}$  are zero).

## Direction of Maximum Variance

- Taking the derivative w.r.t.  $u_1$  and setting to zero gives

$$Su_1 = \lambda_1 u_1$$

- Thus  $u_1$  is an eigenvector of  $S$  (with corresponding eigenvalue  $\lambda_1$ )
- But which of  $S$ 's eigenvectors it is?
- Note that since  $u_1^T u_1 = 1$ , the variance of projected data is

$$u_1^T S u_1 = \lambda_1$$

- Var. is maximized when  $u_1$  is the top eigenvector with largest eigenvalue
- The top eigenvector  $u_1$  is also known as the first Principle Component (PC)
- Other directions can also be found likewise (with each being orthogonal to all previous ones) using the eigendecomposition of  $S$  (this is PCA)

# Steps in Principle Component Analysis

- Center the data (subtract the mean  $\mu = \frac{1}{N} \sum_{i=1}^N x^{(i)}$  from each data point)
- Compute the covariance matrix

$$S = \frac{1}{N} \sum_{i=1}^N x^{(i)} x^{(i)T} = \frac{1}{N} X X^T$$

- Do an eigendecomposition of the covariance matrix  $S$
- Take first  $K$  leading eigenvectors  $\{u_l\}_{l=1, \dots, K}$  with eigenvalues  $\{\lambda_l\}_{l=1, \dots, K}$
- The final  $K$  dim. projection of data is given by

$$Z = U^T X$$

where  $U$  is  $D \times K$  and  $Z$  is  $K \times N$

## PCA as Minimizing the Reconstruction Error

- Assume complete orthonormal basis vector  $u_1, u_2, \dots, u_D$ , each  $u_l \in \mathbb{R}^N$
- We can represent each data point  $x^{(i)} \in \mathbb{R}^D$  exactly using the new basis

$$x^{(i)} = \sum_{l=1}^D z_l^{(i)} u_l$$

$$\begin{bmatrix} x_1^{(i)} \\ x_2^{(i)} \\ \vdots \\ x_D^{(i)} \end{bmatrix} = \begin{bmatrix} u_1 & u_2 & \cdots & u_D \end{bmatrix} * \begin{bmatrix} z_1^{(i)} \\ z_2^{(i)} \\ \vdots \\ z_D^{(i)} \end{bmatrix}$$

- Denoting  $z^{(i)} = [z_1^{(i)} \cdots z_D^{(i)}]^T$ ,  $U = [u_1 \cdots u_D]$ , and using  $U^T U = I$

$$x^{(i)} = U z^{(i)} \quad \text{and} \quad z^{(i)} = U^T x^{(i)}$$

- Also note that each component of vector  $z^{(i)}$  is  $z_l^{(i)} = u_l^T x^{(i)}$

# Reconstruction of Data from Projections

- Reconstruction of  $x^{(i)}$  from  $z^{(i)}$  will be exact if we use all  $D$  basis vectors
- Will be approximate if we only use  $K < D$  basis vectors:

$$x^{(i)} \approx \sum_{l=1}^K z_l^{(i)} u_l$$

- Let's use  $K = 1$  basis vector. Then, the one-dim embedding of  $x^{(i)}$  is

$$z^{(i)} = u_1^T x^{(i)} \quad (z^{(i)} \in \mathbb{R})$$

- We can now try to “reconstruct”  $x^{(i)}$  from its embedding  $z^{(i)}$  as follows

$$\tilde{x}^{(i)} = u_1 z^{(i)} = u_1 u_1^T x^{(i)}$$

- Total error or “loss” in reconstructing all the data points

$$\ell(u_1) = \sum_{i=1}^N \|x^{(i)} - \tilde{x}^{(i)}\|^2 = \sum_{i=1}^N \|x^{(i)} - u_1 u_1^T x^{(i)}\|^2$$

## Direction with Best Reconstruction

- We want to find  $u_1$  that minimize the reconstruction error

$$\ell(u_1) = \sum_{i=1}^N \|x^{(i)} - u_1 u_1^T x^{(i)}\|^2 = \sum_{i=1}^N \left( -u_1^T x^{(i)} (x^{(i)})^T u_1 + (x^{(i)})^T x^{(i)} \right)$$

by using  $u_1^T u_1 = 1$

- Minimizing the error of reconstructing all the data points is equivalent to

$$\max_{u_1: \|u_1\|^2=1} u_1^T \left( \sum_{n=1}^N x^{(i)} (x^{(i)})^T \right) u_1 = \max_{u_1: \|u_1\|^2=1} u_1^T S u_1$$

where  $S$  is the covariance matrix of the data (which are assumed to be centered)

- It is the same objective that we had when we maximized the variance



# Thanks!

Q & A